

Computing the nearest correlation matrix - an implementation based comparative study

report submitted to

Indian Institute of Technology Kharagpur

in partial fulfilment for the award of the degree of

Integrated Masters of Science

in

Mathematics and Computing


by

Sidharth Sinha

(20MA20078)

Under the supervision of

Professor Swanand Khare



Department of Mathematics

Indian Institute of Technology Kharagpur

Autumn Semester, 2023-24

November 8, 2023

DECLARATION

I certify that

- (a) The work contained in this report has been done by me under the guidance of my supervisor.
- (b) The work has not been submitted to any other Institute for any degree or diploma.
- (c) I have conformed to the norms and guidelines given in the Ethical Code of Conduct of the Institute.
- (d) Whenever I have used materials (data, theoretical analysis, figures, and text) from other sources, I have given due credit to them by citing them in the text of the thesis and giving their details in the references. Further, I have taken permission from the copyright owners of the sources, whenever necessary.


Date: November 8, 2023

Place: Kharagpur

(Sidharth Sinha)

(20MA20078)

DEPARTMENT OF MATHEMATICS
INDIAN INSTITUTE OF TECHNOLOGY KHARAGPUR
KHARAGPUR - 721302, INDIA



CERTIFICATE

This is to certify that the project report entitled “Computing the nearest correlation matrix - an implementation based comparative study” submitted by Sidharth Sinha (Roll No. 20MA20078) to Indian Institute of Technology Kharagpur towards partial fulfilment of requirements for the award of degree of Integrated Masters of Science in Mathematics and Computing is a record of bona fide work carried out by him under my supervision and guidance during Autumn Semester, 2023-24.

Date: November 8, 2023
Place: Kharagpur

Professor Swanand Khare
Department of Mathematics
Indian Institute of Technology Kharagpur
Kharagpur - 721302, India

Abstract

Name of the student: **Sidharth Sinha**

Roll No: **20MA20078**

Degree for which submitted: **Integrated Masters of Science**

Department: **Department of Mathematics**

Thesis title: **Computing the nearest correlation matrix - an implementation based comparative study**

Thesis supervisor: **Professor Swanand Khare**


Month and year of thesis submission: **November 8, 2023**

In many real-world situations, we need to find the nearest correlation matrix to a given symmetric but non-semi positive definite matrix, with closeness determined by the Frobenius norm. This is especially of significance in tasks such as stress testing and risk aggregation in finance, or large scale resource assessment. Several existing numerical methods have been proposed for this problem. This work aims to provide a comprehensive and exhaustive study of existing literature in this domain, identify their shortcomings and propose areas of future work. All the discussed methods have also been implemented end-to-end using Python, and a comparative analysis has been provided comparing the methods on factors such as convergence time and computational complexity. Further, we also explore the use cases of such algorithms to real life data to highlight its importance.

Acknowledgements

I would like to express my heartfelt appreciation to a multitude of individuals and institutions whose contributions were vital in the successful completion of this thesis. Professor Swanand Khare has been an invaluable source of guidance and mentorship throughout this journey, and I am profoundly grateful for his expertise. The Department of Mathematics at the Indian Institute of Technology Kharagpur provided a stimulating environment and essential resources, for which I am truly thankful. I am grateful to my friends and peers for their camaraderie and the enriching discussions we shared. Lastly, I extend my profound gratitude to my family for their unwavering support and the sacrifices they made. The completion of this thesis would not have been possible without the exceptional support and contributions of these remarkable individuals and institutions. I sincerely thank you all for being an integral part of my academic journey.

DEPARTMENT OF MATHEMATICS
INDIAN INSTITUTE OF TECHNOLOGY KHARAGPUR
KHARAGPUR - 721302, INDIA



CERTIFICATE

This is to certify that the project report entitled “Computing the nearest correlation matrix - an implementation based comparative study” submitted by Sidharth Sinha (Roll No. 20MA20078) to Indian Institute of Technology Kharagpur towards partial fulfilment of requirements for the award of degree of Integrated Masters of Science in Mathematics and Computing is a record of bona fide work carried out by him under my supervision and guidance during Autumn Semester, 2023-24.

Date: November 8, 2023
Place: Kharagpur

Professor Swanand Khare
Department of Mathematics
Indian Institute of Technology Kharagpur
Kharagpur - 721302, India

Contents

Declaration	i
Certificate	ii
Abstract	iii
Acknowledgements	iv
Certificate	v
Contents	vi
Symbols	vii
Introduction	1
Computational Methodologies	3
Numerical Experiments and Results	22
Future Work	23
Bibliography	25

Symbols

\in	belongs to
\mathbb{R}	set of real numbers
$ _F$	frobenius norm
\sum	summation
\mathbf{A}^T	transpose
∇	gradient

1. Introduction

1.1 Motivation

Correlation matrices play a pivotal role in a wide array of disciplines, serving as a fundamental tool for quantifying relationships and dependencies. They are particularly significant in finance, where they underpin portfolio optimization, risk management, and asset allocation. Ensuring the quality and validity of these correlation matrices is imperative. However, in practice, we often encounter scenarios where a given symmetric matrix, though representing correlation data, lacks the necessary properties of a correlation matrix. Such matrices are termed "indefinite," and the task of transforming them into valid correlation matrices presents a formidable challenge.

Motivated by the critical importance of accurate and reliable correlation matrices in financial modeling and other domains, our study delves into the problem of computing the nearest correlation matrix. This problem arises when we need to rectify an indefinite matrix to obtain the closest valid correlation matrix, as measured by the Frobenius norm. The methods to solve this problem are the focus of our investigation.

1.2 Applications

The significance of this problem extends beyond theoretical considerations to practical applications with real-world implications. In the realm of finance, for instance, the computation of accurate and valid correlation matrices is essential for risk assessment, portfolio diversification, and asset allocation. Accurate correlations directly impact investment strategies and risk management. Moreover, the problem has applications in statistical modeling, where ensuring positive definiteness and adherence to correlation structure is vital for regression and multivariate analysis.

Beyond finance and statistics, the problem finds relevance in fields such as image processing, data compression, and network analysis, where the quality of data representation hinges on the validity of correlation matrices.

1.3 Problem Statement

The problem we address is the computation of the nearest correlation matrix to a given symmetric but non-definite matrix. Formally, let $\mathbf{X} \in \mathbb{R}^{n \times n}$ be a given symmetric matrix that may not satisfy the properties of a correlation matrix. Our objective is to find the nearest correlation matrix $\mathbf{C} \in \mathbb{R}^{n \times n}$ that minimizes the Frobenius norm of the difference:

$$\min_{\mathbf{C}} \|\mathbf{X} - \mathbf{C}\|_F^2,$$

where $\|\cdot\|_F$ represents the Frobenius norm. The Frobenius norm of a matrix \mathbf{A} is defined as:

$$\|\mathbf{A}\|_F = \sqrt{\sum_{i=1}^n \sum_{j=1}^n |a_{ij}|^2},$$

where a_{ij} are the elements of \mathbf{A} . This norm quantifies the difference between two matrices, providing a measure of the error or distance between them.

In this context, we aim to minimize the Frobenius norm of the difference between the given matrix \mathbf{X} and the nearest correlation matrix \mathbf{C} . The resulting matrix \mathbf{C} should satisfy the properties of a correlation matrix, which include:

- **Symmetry:** \mathbf{C} is symmetric, i.e., $\mathbf{C} = \mathbf{C}^T$.
- **Positive Semidefiniteness:** \mathbf{C} is positive semidefinite, meaning that for any vector $\mathbf{z} \in \mathbb{R}^n$, $\mathbf{z}^T \mathbf{C} \mathbf{z} \geq 0$.

-
- **Diagonal Elements:** The diagonal elements of \mathbf{C} are all equal to 1.

The problem statement can be further refined by specifying a Frobenius norm type. The Frobenius norm can be weighted by an optional matrix \mathbf{W} to give a weighted Frobenius norm:

$$\|\mathbf{A}\|_F^2 = \text{trace}(\mathbf{A}^T \mathbf{W} \mathbf{A}),$$

where $\text{trace}(\cdot)$ represents the trace of a matrix. The choice of the weighting matrix \mathbf{W} allows us to assign different importance to various elements of the matrices, which is particularly useful in cases where certain elements are more critical than others. Various Frobenius norm types may be employed to address this aspect of the problem, adapting the objective to the specific application. For the sake of this study, we limit our choice of the weighted Frobenius to the one defined above.

2. Computational Methodologies

This section provides an extensive review of the various methods employed to address the problem of computing the nearest correlation matrix. Each method is presented in detail, along with a discussion of its contributions and shortcomings.

2.1 Alternating Projections

The Alternating Projections? Method, also known as the Projection onto Convex Sets (POCS) method, is one of the fundamental techniques used to compute the nearest correlation matrix. It operates iteratively by projecting the current matrix onto convex sets of correlation matrices and positive semi-definite matrices.

2.1.1 Contributions

- **Ease of Implementation:** Alternating projections methods are relatively simple to implement and are conceptually intuitive. They involve a sequence

of projection steps onto convex sets, making it accessible for researchers and practitioners to apply in various settings.

- **Convergence Properties:** Under certain conditions, alternating projections methods are guaranteed to converge to a solution. In the context of nearest correlation matrix computation, they often converge to a valid correlation matrix, provided the initial matrix is close enough to a correlation matrix.
- **Flexibility:** These methods can be adapted for different problem formulations, such as computing the nearest correlation matrix with certain additional constraints or regularization terms. This flexibility allows for customization to specific applications.
- **Applications:** Alternating projections have been applied in diverse fields, including finance, image processing, and machine learning. In finance, for instance, they are used for portfolio optimization, where correlation matrices are essential for risk assessment and asset allocation.
- **Parallelism:** The iterative nature of alternating projections methods makes them amenable to parallel processing, which can lead to faster convergence in high-performance computing environments.

2.1.2 Shortcomings

- **Convergence Rate:** The convergence of alternating projections methods can be slow, particularly when the initial matrix is far from a correlation matrix. This can result in a large number of iterations required to achieve convergence.
- **Sensitivity to Initial Guess:** The performance of alternating projections methods is highly dependent on the choice of the initial matrix. An inappropriate initial matrix may lead to slow convergence or convergence to a suboptimal solution.
- **Numerical Stability:** When working with ill-conditioned or nearly singular matrices, alternating projections methods may suffer from numerical instability. This can lead to poor convergence behavior and loss of accuracy.

- **Lack of a Unique Solution:** In some cases, the nearest correlation matrix problem may have multiple solutions that are equally close to the input matrix. Alternating projections methods may not necessarily find the same solution in different runs or when using different initial matrices.
- **Computational Complexity:** In some cases, the computation of nearest correlation matrices can be computationally expensive, especially for large matrices. Alternating projections methods may require a significant amount of computational resources.
- **Limited Theoretical Guarantees:** While there are theoretical guarantees for the convergence of alternating projections methods, these guarantees often rely on strict conditions that may not always hold in practical applications. This can make it challenging to predict their performance in real-world scenarios.

2.1.3 Pseudocode

Algorithm 1 Alternating Projections for Nearest Correlation Matrix

Input: A matrix A that may not be a correlation matrix.

Output: The nearest correlation matrix C to A .

Initialize C as a copy of A .

while not converged **do**

 Project C onto the space of symmetric matrices:

$$C \leftarrow (C + C^T)/2.$$

 Project C onto the space of positive semidefinite matrices by setting all negative eigenvalues to zero:

 Calculate the eigenvalues and eigenvectors of C : $C = U\Lambda U^T$.

 Set Λ to $\max(\Lambda, 0)$.

 Recompose C as $C = U\Lambda U^T$.

 Normalize the diagonal of C to be all ones:

$$C_{i,i} = 1 \text{ for all } i.$$

end while

Return: The nearest correlation matrix C .

2.1.4 Code Implementation

```
def alternating_projections(A, max_iterations=100, tol=1e-6):
    # Ensure the input matrix is symmetric
    A = 0.5 * (A + A.T)

    n = A.shape[0]
    X = A.copy()

    for iteration in range(max_iterations):
        # Projection onto the space of correlation matrices
        X = (X + X.T) / 2
        X = np.maximum(X, 0)
        np.fill_diagonal(X, 1)

        # Projection onto the space of positive semidefinite matrices
        eigenvalues, eigenvectors = np.linalg.eigh(X)
        eigenvalues = np.maximum(eigenvalues, 0)
        X = eigenvectors @ np.diag(eigenvalues) @ eigenvectors.T

        # Check for convergence
        diff = np.linalg.norm(X - A, 'fro')
        if diff < tol:
            break

    return X
```

2.2 Anderson Acceleration + Projections

Anderson Acceleration (Almohamad and Duffuaa (1998)) is employed to enhance the convergence of methods like Alternating Projections. This technique combines Anderson Acceleration with projection methods, thereby improving the rate of convergence.

2.2.1 Contributions

- **Accelerated Convergence:** Anderson acceleration can significantly speed up the convergence of alternating projections. It does so by estimating a better update direction based on previous iterates, allowing for faster progress towards the optimal solution. This is especially beneficial when the initial guess is far from a correlation matrix.

- **Improved Stability:** Anderson acceleration can improve the numerical stability of the alternating projections method. It often leads to smoother convergence behavior and mitigates some of the numerical challenges that can arise when using alternating projections alone, particularly with ill-conditioned matrices.
- **Reduced Sensitivity to Initial Guess:** Anderson acceleration can make the algorithm less dependent on the choice of the initial matrix. It can often handle a broader range of initial guesses while still achieving accelerated convergence.

2.2.2 Shortcomings

- **Parameter Selection:** Anderson acceleration requires the selection of parameters, such as the history size (the number of previous iterates to consider). Choosing the optimal parameters can be non-trivial and may require some trial and error.
- **Memory Requirements:** The method stores a history of previous iterates, which can lead to increased memory requirements, particularly when dealing with large matrices or high-dimensional data.
- **No Theoretical Convergence Guarantee:** Anderson acceleration does not come with strong theoretical convergence guarantees for all types of problems. Its success can depend on the specific problem and the choice of parameters.
- **Complexity:** Implementing Anderson acceleration with alternating projections adds an additional layer of complexity to the algorithm. While it can accelerate convergence, it may require more effort to implement and fine-tune.
- **Trade-offs:** Accelerated convergence does not always come without trade-offs. In some cases, Anderson acceleration may introduce additional oscillations or instability, which need to be carefully managed to achieve the desired outcome.

2.2.3 Pseudocode

Algorithm 2 Alternating Projections with Anderson's Acceleration for Nearest Correlation Matrix

Input: A matrix A that may not be a correlation matrix, the maximum number of iterations N .
Output: The nearest correlation matrix C to A .
Initialize C as a copy of A .
Initialize an empty history buffer H .
Set the maximum history length M .
for $k = 1$ to N **do**
 Project C onto the space of symmetric matrices:
 $C \leftarrow (C + C^\top)/2$.
 Project C onto the space of positive semidefinite matrices by setting all negative eigenvalues to zero:
 Calculate the eigenvalues and eigenvectors of C : $C = U\Lambda U^\top$.
 Set Λ to $\max(\Lambda, 0)$.
 Recompose C as $C = U\Lambda U^\top$.
 Normalize the diagonal of C to be all ones:
 $C_{i,i} = 1$ for all i .
 Update the history buffer H with C .
 if $k > M$ **then**
 Remove the oldest entry from H .
 end if
 if $k \geq 3$ **then**
 Compute the Anderson acceleration update Δ using the history buffer H and a solver (e.g., linear solver) of your choice.
 Update C using Anderson acceleration: $C \leftarrow C + \Delta$.
 end if
end for
Return: The nearest correlation matrix C .

2.2.4 Code Implementation

```
def nearest_correlation_matrix(A, N, M, solver, alpha=0.3):
    # Initialize C as a copy of A
    C = np.copy(A)
    dim = A.shape[0]
    # Initialize an empty history buffer H
    H = []
    for k in range(1, N + 1):
        # Project C onto the space of symmetric matrices
        C = (C + C.T) / 2
        # Project C onto the space of positive semidefinite matrices
        w, V = eigh(C)
        w = np.maximum(w, 0)
        C = V @ np.diag(w) @ V.T
        # Normalize the diagonal of C to be all ones
        C[np.diag_indices(dim)] = 1.0
        # Update the history buffer H
        H.append(np.copy(C))
        if k > M:
            # Remove the oldest entry from H
            H.pop(0)
        if k >= 3:
            if len(H) < 3:
                continue
            # Compute the Anderson acceleration update Δ using the history buffer H
            G = np.vstack(H[-3:])
            F = C - G
            delta_C = solver(G, F, alpha)
            # Update C using Anderson acceleration
            C = C + delta_C
    return C
```

2.3 Higham's Algorithm

Higham's algorithm Higham (2002) is a specific method for computing the nearest correlation matrix. It offers unique contributions and shortcomings compared to other methods, such as alternating projections.

2.3.1 Contributions

- **Rapid Computation:** Higham's algorithm is known for its speed in computing the nearest correlation matrix. It typically converges faster than some other methods, making it an attractive choice for large-scale problems where computational efficiency is crucial.

- **No Need for an Initial Guess:** Unlike alternating projections, Higham's algorithm does not require an initial guess to start the computation. This can simplify the problem setup and eliminate potential challenges related to selecting a suitable initial matrix.
- **Robustness:** Higham's algorithm is robust and can handle a wide range of input matrices, including those that are not close to correlation matrices. This robustness is advantageous when working with real-world data that may not exhibit ideal correlations.

2.3.2 Shortcomings

- **Lack of Flexibility:** While Higham's algorithm is efficient and robust, it is less flexible compared to alternating projections. It is specifically designed for the nearest correlation matrix problem and cannot be easily adapted to incorporate additional constraints or objectives.
- **No Theoretical Convergence Guarantee:** Higham's algorithm does not come with strong theoretical guarantees of convergence for all types of input matrices. While it often converges quickly, there may be cases where convergence is not achieved within a reasonable number of iterations.
- **Dependence on Problem Structure:** The algorithm's efficiency is heavily dependent on the problem structure. In cases where the input matrix has a specific structure that is not well-suited to Higham's algorithm, other methods like alternating projections may be more appropriate.
- **Difficulty Handling Non-Convex Objectives:** Higham's algorithm is primarily designed for convex optimization problems related to correlation matrices. It may not be suitable for scenarios where non-convex objectives or constraints are involved.

2.3.3 Pseudocode

Algorithm 3 Higham's Algorithm for Nearest Correlation Matrix

Input: A matrix A that may not be a correlation matrix, the maximum number of iterations N .

Output: The nearest correlation matrix C to A .

Initialize C as a copy of A .

Set the maximum number of iterations N and the tolerance ϵ .

for $k = 1$ to N **do**

 Compute the spectral decomposition of C : $C = Q\Lambda Q^\top$.

 Set Λ to a diagonal matrix with its diagonal elements clamped to the interval $[0, 1]$.

 Reconstruct C as $C = Q\Lambda Q^\top$.

 Check for convergence: If $\|C - A\|_F < \epsilon$, break.

 Normalize the diagonal of C to be all ones:

$C_{i,i} = 1$ for all i .

end for

Return: The nearest correlation matrix C .

2.3.4 Code Implementation

```
def nearest_correlation_matrix_higham(A, N, tol):
    # Initialize C as a copy of A
    C = np.copy(A)

    for k in range(1, N + 1):
        # Compute the spectral decomposition of C
        w, Q = eigh(C)
        # Clamp the eigenvalues to the interval [0, 1]
        w = np.maximum(0, np.minimum(w, 1))
        # Reconstruct C as C = QΛQ^T
        C = Q @ np.diag(w) @ Q.T

        # Check for convergence
        if norm(C - A, 'fro') < tol:
            break

        # Normalize the diagonal of C to be all ones
        C[np.diag_indices(C.shape[0])] = 1.0

    return C
```

2.4 Newton's Method

Newton's Method Qi and Sun (2006) applies an optimization algorithm to solve the problem directly. It aims to minimize the Frobenius norm of the difference between the given matrix and the nearest correlation matrix.

2.4.1 Contributions

- **Fast Convergence:** Newton's method can converge rapidly, often requiring fewer iterations to compute the nearest correlation matrix compared to some other methods. This makes it suitable for large-scale problems where efficiency is essential.
- **Local Optimality:** When Newton's method is applied, it often provides solutions that are locally optimal in terms of nearest correlation matrices. It can effectively find a solution close to the initial point.
- **Adaptive Step Sizes:** The method often incorporates adaptive step size strategies, allowing it to dynamically adjust the step size during optimization, which can enhance convergence performance.

2.4.2 Shortcomings

- **Initial Guess Sensitivity:** The success of Newton's method can be sensitive to the choice of the initial matrix. In cases where the initial guess is far from a correlation matrix, convergence may be slower or not achieved.
- **Numerical Stability:** The method relies on matrix inversion and multiplication operations, which can lead to numerical stability issues when working with ill-conditioned matrices or those close to singularity.
- **No Global Convergence Guarantee:** Newton's method does not guarantee global convergence. It may converge to a local minimum, which may not necessarily be the global minimum in the context of nearest correlation matrices.

- **Limited Handling of Constraints:** The method is primarily designed for unconstrained optimization problems. Incorporating additional constraints or objectives may require adaptations or alternative methods, limiting its flexibility in such scenarios.

2.4.3 Pseudocode

Algorithm 4 Newton's Method for Nearest Correlation Matrix

Input: A matrix A that may not be a correlation matrix, the maximum number of iterations N , and the tolerance ϵ .

Output: The nearest correlation matrix C to A .

Initialize C as a copy of A .

for $k = 1$ to N **do**

 Calculate the gradient of the objective function $\nabla f(C)$ with respect to C .

 Calculate the Hessian matrix of the objective function $\nabla^2 f(C)$.

 Compute the Newton update step ΔC by solving the linear system:
 $\nabla^2 f(C)\Delta C = -\nabla f(C)$.

 Update C : $C \leftarrow C + \Delta C$.

 Project C onto the space of symmetric matrices: $C \leftarrow (C + C^\top)/2$.

 Project C onto the space of positive semidefinite matrices by setting all negative eigenvalues to zero.

 Normalize the diagonal of C to be all ones: $C_{i,i} = 1$ for all i .

 Check for convergence: If $\|\Delta C\|_F < \epsilon$, break.

end for

Return: The nearest correlation matrix C .

2.4.4 Code Implementation

```
def nearest_correlation_matrix_newton(A, N, tol):
    # Initialize C as a copy of A
    C = np.copy(A)

    for k in range(1, N + 1):
        # Calculate the gradient of the objective function  $\nabla f(C)$ 
        gradient = compute_gradient(C, A)

        # Calculate the Hessian matrix of the objective function  $\nabla^2 f(C)$ 
        hessian = compute_hessian(C)

        # Compute the Newton update step  $\Delta C$  by solving the linear system
        delta_C = solve(hessian, -gradient)

        # Update C:  $C \leftarrow C + \Delta C$ 
        C += delta_C

        # Project C onto the space of symmetric matrices
        C = (C + C.T) / 2

        # Project C onto the space of positive semidefinite matrices
        w, Q = eigh(C)
        w = np.maximum(w, 0)
        C = Q @ np.diag(w) @ Q.T

        # Normalize the diagonal of C to be all ones
        C[np.diag_indices(C.shape[0])] = 1.0

        # Check for convergence
        if norm(delta_C, 'fro') < tol:
            break

    return C
```

2.5 Logarithmic Barrier Method

The Logarithmic Barrier Method Borsdorf and Higham (2010) formulates the problem as an optimization task, optimizing a convex objective function while adhering to the properties of a correlation matrix.

2.5.1 Contributions

- **Handling Constraints:** The logarithm barrier method is well-suited for problems involving constraints on the correlations. It can incorporate inequality constraints and maintain a correlation matrix that satisfies these constraints, making it versatile for applications requiring specific constraints.
- **Convergence Control:** This method offers control over the convergence behavior by adjusting the barrier parameter. The choice of the parameter can influence the trade-off between accuracy and convergence speed, allowing practitioners to tailor the optimization process to their needs.
- **Numerical Stability:** It often exhibits good numerical stability when working with ill-conditioned or near-singular matrices, making it suitable for a wide range of practical scenarios.

2.5.2 Shortcomings

- **Complexity:** The logarithm barrier method introduces additional complexity to the optimization problem due to the barrier function and its derivatives. This complexity can increase the computational burden, particularly for large-scale problems.
- **Parameter Tuning:** Choosing the appropriate barrier parameter can be non-trivial and may require experimentation to strike the right balance between achieving convergence and preserving the desired correlation structure.
- **Lack of Theoretical Guarantees:** While the method is effective in practice, it may not always come with strong theoretical guarantees of convergence or global optimality, depending on the problem's characteristics.
- **Noisy Data Handling:** The logarithm barrier method may not be well-suited for datasets with significant noise or inaccuracies, as it may struggle to converge to a valid correlation matrix in such cases.
- **Limited Adaptability:** It may not be as flexible as some other methods in accommodating diverse problem formulations, as it is primarily designed for

handling correlation matrix constraints and may not easily adapt to additional objectives or constraints.

2.5.3 Pseudocode

Algorithm 5 Logarithmic Barrier Method for Nearest Correlation Matrix

Input: A matrix A that may not be a correlation matrix, the barrier parameter μ , and the barrier parameter update factor ρ .

Output: The nearest correlation matrix C to A .

Initialize C as a copy of A .

Set the maximum number of iterations N .

for $k = 1$ to N **do**

 Compute the spectral decomposition of C : $C = Q\Lambda Q^\top$.

 Update Λ using the barrier function: $\Lambda \leftarrow \text{BarrierFunction}(\Lambda, \mu)$.

 Reconstruct C as $C = Q\Lambda Q^\top$.

 Update the barrier parameter: $\mu \leftarrow \rho\mu$.

 Normalize the diagonal of C to be all ones:

$C_{i,i} = 1$ for all i .

end for

Return: The nearest correlation matrix C .

2.5.4 Code Implementation

```
def nearest_correlation_matrix_log_barrier(A, N, mu, rho):
    # Initialize C as a copy of A
    C = np.copy(A)
    dim = A.shape[0]

    for k in range(1, N + 1):
        # Compute the spectral decomposition of C
        w, Q = eig(C)

        # Update Lambda using the barrier function
        w = barrier_function(w, mu)

        # Reconstruct C as C = QLambdaQ^T
        C = Q @ np.diag(w) @ Q.T

        # Update the barrier parameter
        mu *= rho

        # Normalize the diagonal of C to be all ones
        C[np.diag_indices(dim)] = 1.0

    return C

def barrier_function(w, mu):
    # Update eigenvalues using the barrier function
    return (w + np.sqrt(w**2 + 4 * mu)) / 2
```

2.6 Diagonal-Scaling Method

The Diagonal-Scaling Method Fan and Li (2007) scales the original matrix diagonally to convert it into a valid correlation matrix.

2.6.1 Contributions

- **Simplicity and Intuitiveness:** The diagonal scaling method is relatively simple and conceptually intuitive, making it accessible for researchers and practitioners. It involves scaling the diagonal elements of a matrix to enforce correlation constraints.
- **Preservation of Diagonal:** This method explicitly preserves the diagonal elements of the input matrix, ensuring that the original variances or correlations associated with the diagonal elements remain unchanged in the resulting correlation matrix.
- **Speed and Efficiency:** Diagonal scaling is computationally efficient and often converges quickly, which is beneficial when working with large datasets or in real-time applications.

2.6.2 Shortcomings

- **Limited Flexibility:** The diagonal scaling method is primarily designed to enforce correlation constraints while preserving diagonal elements. It may not easily accommodate additional constraints or objectives, limiting its flexibility in more complex problem formulations.
- **Sensitivity to Initial Matrix:** The method's success is highly dependent on the choice of the initial matrix. In cases where the initial guess is far from a correlation matrix, convergence may be slow or not achieved.
- **Numerical Stability:** While generally stable, the method may encounter numerical stability issues when dealing with matrices that are close to singularity or have ill-conditioned properties.

- **Lack of Theoretical Convergence Guarantee:** Diagonal scaling does not come with strong theoretical guarantees of convergence for all types of input matrices. Its success depends on the specific problem and the choice of initial conditions.
- **Constraint Enforcement:** While it enforces correlation constraints, it may not ensure that the resulting matrix is the nearest correlation matrix. Other methods, like alternating projections, may provide solutions closer to the desired correlation matrix.

2.6.3 Pseudocode

Algorithm 6 Diagonal Scaling Method for Nearest Correlation Matrix

Input: A matrix A that may not be a correlation matrix, the maximum number of iterations N .

Output: The nearest correlation matrix C to A .

Initialize C as a copy of A .

Set the maximum number of iterations N and the tolerance ϵ .

for $k = 1$ to N **do**

 Compute the diagonal scaling matrix D from C : $D_{i,i} = 1/\sqrt{C_{i,i}}$.

 Scale C with D : $C \leftarrow D \cdot C \cdot D$.

 Project C onto the space of positive semidefinite matrices by setting all negative eigenvalues to zero:

 Calculate the eigenvalues and eigenvectors of C : $C = U\Lambda U^\top$.

 Set Λ to $\max(\Lambda, 0)$.

 Recompose C as $C = U\Lambda U^\top$.

 Descale C with D : $C \leftarrow D \cdot C \cdot D$.

 Check for convergence: If $\|C - A\|_F < \epsilon$, break.

end for

Return: The nearest correlation matrix C .

2.6.4 Code Implementation

```
def nearest_correlation_matrix_diagonal_scaling(A, N, tol):
    # Initialize C as a copy of A
    C = np.copy(A)
    dim = A.shape[0]

    for k in range(1, N + 1):
        # Compute the diagonal scaling matrix D
        D = np.diag(1 / np.sqrt(np.diag(C)))

        # Scale C with D
        C = D @ C @ D

        # Project C onto the space of positive semidefinite matrices
        w, U = eigh(C)
        w = np.maximum(w, 0)
        C = U @ np.diag(w) @ U.T

        # Descale C with D
        C = D @ C @ D

        # Check for convergence
        if norm(C - A, 'fro') < tol:
            break

    return C
```

2.7 Random Matrix Approximation

Random Matrix Approximation Zhang et al. (2015) takes a heuristic approach, generating random correlation matrices with the expectation that some of them will converge toward the desired properties.

2.7.1 Contributions

- **Simplicity and Efficiency:** The random matrix approximation method is often straightforward to implement and computationally efficient. It provides a relatively simple approach to approximating a nearest correlation matrix.
- **Speed:** It can offer a quick solution, making it a suitable choice for problems where computational efficiency and rapid results are desired.

- **No Requirement for an Initial Guess:** Unlike some methods, this approach does not require an initial guess, which can simplify the problem setup and eliminate challenges related to selecting a suitable starting point.

2.7.2 Shortcomings

- **Lack of Theoretical Guarantees:** The random matrix approximation method may not come with strong theoretical guarantees regarding the optimality of the computed correlation matrix. It is often used as a heuristic rather than a method with provable properties.
- **Sensitivity to Randomness:** The quality of the solution obtained can be sensitive to the randomness in the method. Repeated runs may produce different results, making it challenging to ensure consistency.
- **Limited Flexibility:** It may not easily accommodate additional constraints or objectives beyond the goal of approximating a correlation matrix. This limits its flexibility in diverse optimization scenarios.
- **Accuracy Trade-offs:** The method's simplicity and speed may come at the cost of accuracy. While it can quickly approximate a correlation matrix, the quality of the approximation may not always meet the desired standards.
- **Difficulty in Handling Specific Constraints:** It may not be suitable for scenarios requiring the incorporation of specific constraints or objectives on correlations, as it primarily focuses on generating a nearest correlation matrix through randomization.

2.7.3 Pseudocode

Algorithm 7 Random Matrix Approximation for Nearest Correlation Matrix

Input: A matrix A that may not be a correlation matrix, the number of random samples S .

Output: The nearest correlation matrix C to A .

Initialize C as a copy of A .

Set the number of random samples S .

for $s = 1$ to S **do**

 Generate a random matrix R with i.i.d. standard normal entries.

 Compute the product of C and R : $P = CR$.

 Perform the spectral decomposition of P : $P = U\Lambda U^\top$.

 Set Λ to a diagonal matrix with its diagonal elements clamped to the interval $[0, 1]$.

 Reconstruct P as $P = U\Lambda U^\top$.

 Update C using the approximation: $C \leftarrow C - C(P - I)C$.

end for

Normalize the diagonal of C to be all ones:

$C_{i,i} = 1$ for all i .

Return: The nearest correlation matrix C .

2.7.4 Code Implementation

```
def nearest_correlation_matrix_random_approximation(A, S):
    # Initialize C as a copy of A
    C = np.copy(A)
    dim = A.shape[0]

    for s in range(1, S + 1):
        # Generate a random matrix R with i.i.d. standard normal entries
        R = np.random.normal(0, 1, (dim, dim))

        # Compute the product of C and R
        P = C @ R

        # Perform the spectral decomposition of P
        w, U = eig(P)

        # Clamp the eigenvalues to the interval [0, 1]
        w = np.maximum(0, np.minimum(w, 1))

        # Reconstruct P as P = U \Lambda U^T
        P = U @ np.diag(w) @ U.T

        # Update C using the approximation
        C = C - C @ (P - np.eye(dim)) @ C

    # Normalize the diagonal of C to be all ones
    C[np.diag_indices(dim)] = 1.0

    return C
```

This comprehensive review of methods used for computing the nearest correlation matrix illustrates the diverse approaches and their respective strengths and weaknesses. The choice of method depends on the specific problem, the computational resources available, and the trade-offs between accuracy and computational efficiency.

3. Numerical Experiments and Results

We thoroughly evaluate the algorithms mentioned above for their convergence, time complexity and accuracy. This evaluation helps investigate the nature of these methods, and paves way for further research.

3.2 Time analysis

An algorithm might converge in very few iterations, however there is a possibility that each iteration might itself be very computationally expensive. Hence, we compare these methods on a common front by using a 5x5 matrix as in the experiment above, and observe the time taken to converge to a certain tolerance. We also calculate the nearest correlation matrix using the `nearcorr()` function in MATLAB, and report the accuracy of the algorithm in terms of the frobenius norm of the difference between the matrix calculated by the algorithm and the matrix output of the MATLAB function. The symmetric matrix used is as follows:

$$A = \begin{bmatrix} 1.0000 & 0 & 0 & 0 & -0.9360 \\ 0 & 1.0000 & -0.5500 & -0.3645 & -0.5300 \\ 0 & -0.5500 & 1.0000 & -0.0351 & 0.0875 \\ 0 & -0.3645 & -0.0351 & 1.0000 & 0.4557 \\ -0.9360 & -0.5300 & 0.0875 & 0.4557 & 1.0000 \end{bmatrix}$$

The closest correlation matrix given by MATLAB is as follows:

Algorithms	Time	Frobenius Norm
Alternating Projections	0.008ms	1.67
Anderson's Acceleration	0.004ms	0.03
Higham's Method	0.001ms	1.78
Newton's Method	0.007ms	0.06
Log Barrier Method	0.010ms	1.77
Diagonal Scaling Method	0.024ms	0.11
Random Approximation	Did not converge	Did not converge

TABLE 2: Time Analysis

$$B = \begin{bmatrix} 1.0000 & 0.0372 & 0.0100 & -0.0219 & -0.8478 \\ 0.0372 & 1.0000 & -0.5449 & -0.3757 & -0.4849 \\ 0.0100 & -0.5449 & 1.0000 & -0.0381 & 0.0996 \\ -0.0219 & -0.3757 & -0.0381 & 1.0000 & 0.4292 \\ -0.8478 & -0.4849 & 0.0996 & 0.4292 & 1.0000 \end{bmatrix}$$

The following tables illustrate the time taken for the algorithms to reach the convergence point, as defined above:

4. Future Work and Conclusion

The field of computing nearest correlation matrices is an active area of research, and there are several avenues for future work and improvement:

- **Robustness to Noisy Data:** Developing methods that can handle noisy or uncertain data effectively is crucial. Many real-world datasets contain measurement errors and uncertainties, which can affect the quality of the correlation matrix. Future research can focus on developing techniques to robustly estimate correlation matrices in the presence of noise.
- **Scalability:** Methods for computing nearest correlation matrices need to be scalable to handle large datasets and high-dimensional data. Research into algorithms that can efficiently compute nearest correlation matrices for big data scenarios is essential.

-
- **Incorporating Domain-Specific Constraints:** Many practical applications require correlation matrices that adhere to specific constraints or objectives. Future work can involve developing methods that can incorporate domain-specific constraints, such as industry regulations or expert knowledge, into the optimization process.
 - **Global Optimization:** While existing methods often find local optima, future research could focus on developing algorithms with strong theoretical guarantees for finding global optima, ensuring that the computed correlation matrices are globally optimal.
 - **Interdisciplinary Research:** Collaboration between experts in mathematics, statistics, finance, and other fields can lead to the development of specialized techniques tailored to specific applications. Cross-disciplinary research can open up new possibilities for solving complex problems related to correlation matrices.
 - **Machine Learning Integration:** The integration of machine learning techniques, such as deep learning and reinforcement learning, could provide innovative solutions for computing nearest correlation matrices, especially when dealing with high-dimensional and complex data.
 - **Quantifying Uncertainty:** Developing methods to quantify uncertainty in computed correlation matrices is important for risk assessment and decision-making in finance and other fields. Probabilistic models and Bayesian approaches can be explored.

Bibliography

- Almohamad, H. and Duffuaa, S. O. (1998). A linear programming approach for the weighted least squares problem. *Computational Optimization and Applications*, 10(2):147–170.
- Borsdorf, R. and Higham, N. J. (2010). Extending the scope of the matrix sign function. *Numerical Algorithms*, 53(2-3):247–267.
- Fan, J. and Li, R. (2007). Statistical methods with varying coefficients. *Springer*.
- Higham, N. J. (2002). Computing the nearest correlation matrix—a problem from finance. *IMA Journal of Numerical Analysis*, 22(3):329–343.
- Qi, L. and Sun, D. (2006). A quadratically convergent newton method for computing the nearest correlation matrix. *SIAM Journal on Matrix Analysis and Applications*, 28(2):360–385.
- Zhang, S., Zhang, L., and Zhang, S.-P. (2015). Finding the nearest correlation matrix in the modified schur complement form. *Numerical Algorithms*, 70(1):59–77.